

Final Project

Ruben Solis

Due: 29 April 2020

We begin this project by bringing the data into R. The data comes from the kernlab package, and we seek to use several classification methods to determine if an email is a regular email or spam email.


```
library(kernlab)
data(spam)
dat <- spam
```

The data set contains 4601 emails in total and of which 1813 were spam emails, so approximately 40% of the emails in the data set are spam. Looking at the data we see that the type of each variable, except for our target variable, is continuous. Our target variable type is a character variable that is recorded as either spam or nonspam. Moreover, we note that there are no missing values in our dataset.

We next split the data into a training and test set with the following code

```
set.seed(5474)
training.data.index <- sample(1:nrow(dat), 0.667*nrow(dat))
training <- dat[training.data.index, ]
test <- dat[-training.data.index, ]


library(cluster)
library(factoextra)
distance <- get_dist(dat[, -58], method = "euclidean")
fviz_nbclust(dat[, -58], hcut, k.max=5) + geom_vline(xintercept = 2, linetype = 2)
```



The figure above shows that the optimal number of clusters needed is 2, as we suspected.

We will use the distance matrix to build the dendrogram. Furthermore, we will use Ward's method.

```
hc.out <- hclust(distance, method = "ward.D2")
hc.clusters <- cutree(hc.out, k=2)
plot(hc.out)
rect.hclust(hc.out, k=2, border=c("red", "green"))
```



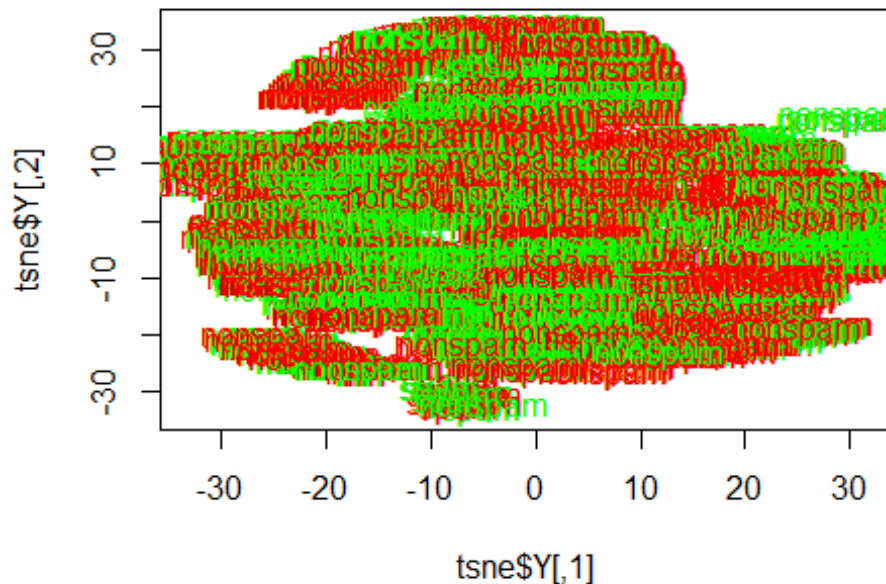
```
table(hc.clusters)/4601 * 100
```

```
## hc.clusters
##      1      2
## 87.7418 12.2582
```

We see that the clustering did not produce great results. The original data set had approximately 40% spam emails whereas here we have about only 12%. We now move on to tSNE.

```
library(Rtsne)
labs <- dat[, 58]
tsne <- Rtsne(dat[, -58], dims=2, perplexity=30, check_duplicates = FALSE,
max_iter=500)
plot(tsne$Y, t='n', main = "t-distributed Stochastic Neighbour Embedding")
text(tsne$Y, labels = labs, col = c("green", "red"))
```

t-distributed Stochastic Neighbour Embedding



We see that the tsne did not really perform that well.

We next fit models using three different statistical learning methods to classify an email as either spam or nonspam; these methods are listed in each of the subsection below.

Our first method is linear discriminate analysis (LDA). The code below is used to fit said model.

```
library(MASS)
spam.lda <- lda(type ~., data = training, CV = FALSE)
spam.lda.pred <- as.vector(predict(spam.lda, test)$x)
```

Our second method is a logistic regression model selected via backward elimination based on BIC. The code below is used to fit the model.

```
fit.full <- glm(type~., data = training, family = binomial)
spam.back <- step(fit.full, direction = "backward",
                  k = log(nrow(dat)), trace = FALSE)
spam.back.pred <- predict(spam.back, newdata = test, type = "response")
```

```
library(rpart)
library(RColorBrewer)
library(rJava)
library(partykit)
library(rpart.plot)
control0 <- rpart.control(minsplit=10, minbucket=3, maxdepth=10, cp=0.01,
                           maxcompete=4, maxsurrogate=3, usesurrogate=2,
                           surrogatestyle=0, xval=10)
```

```
tre0 <- rpart(type ~., data = training, method="class", control=control0,
             parms=list(split="information"), model = T)
```

Next, we need to prune the tree and find the optimal one. The code below will accomplish this task.

```
cv.error <- (tre0$cptable)[,4]
a0 <- 1      # IF a0=0, THEN 0SE
SE1 <- min(cv.error) + a0*((tre0$cptable)[,5])[which.min(cv.error)]
position <- min((1:length(cv.error))[cv.error <= SE1])
n.size <- (tre0$cptable)[,2] + 1 # TREE SIZE IS ONE PLUS NUMBER OF SPLITS.
best.size <- n.size[position]; best.size

## 7
## 9

best.cp <- sqrt(tre0$cptable[position,1] * tre0$cptable[(position-1),1])
best.tree <- prune(tre0, cp=best.cp)
pred <- predict(best.tree, newdata =test, type="prob", na.action = na.pass)
spam.tree.pred <- pred[,1]
```

Our next classification model is a random forest. We have elected to use the default values as arguments/parameters when fitting this model. The code for this is below.

```
library(randomForest, quietly = TRUE)
set.seed(5474)
spam.rf <- randomForest(type~., data = training)
spam.rf.pred <- as.data.frame(predict(spam.rf, test, type = "prob"))$spam
```

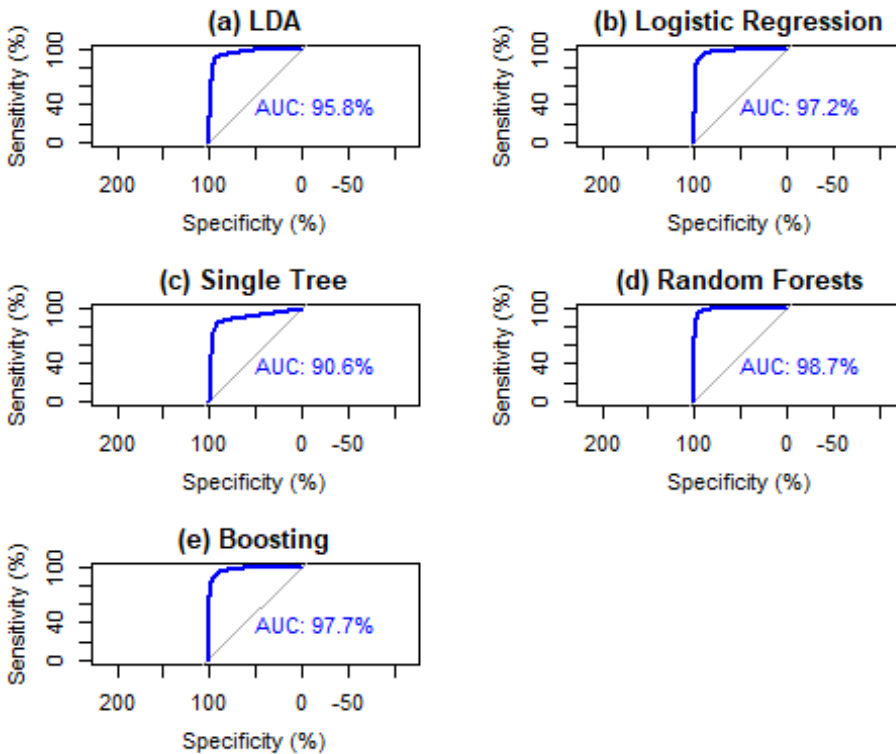
The goal is to make a weaker learner into stronger learner.

```
library(ada)
stump <- rpart.control(cp=-1, maxdepth=1, minsplit=0)
fit.stump <- ada(type ~., data=training, iter=500, loss="e",
               type="discrete", control=stump);
fit1.stump <- addtest(x=fit.stump, test.x = test[, -58], test.y = test[, 58])
spam.boost.pred <- predict(fit.stump, newdata=test, type="probs")[, 2]
```

Finally, we evaluate the performance of these models by comparing their area under the receiver operating characteristic curve (AUC) values.

```
library(pROC, quietly = TRUE)
par(mfrow=c(3,2), mar=rep(4,4))
plot.roc(test$type, spam.lda.pred, main="(a) LDA", percent=TRUE,
         print.auc=TRUE, print.auc.cex=1.0, col="blue")
plot.roc(test$type, spam.back.pred, main="(b) Logistic Regression",
         percent=TRUE,
         print.auc=TRUE, print.auc.cex=1.0, col="blue")
plot.roc(test$type, spam.tree.pred, main="(c) Single Tree", percent=TRUE,
         print.auc=TRUE, print.auc.cex=1.0, col="blue")
plot.roc(test$type, spam.rf.pred, main="(d) Random Forests", percent=TRUE,
         print.auc=TRUE, print.auc.cex=1.0, col="blue")
```

```
plot.roc(test$type, spam.boost.pred, main = "(e) Boosting", percent=TRUE,
         print.auc=TRUE, print.auc.cex=1.0, col="blue")
```

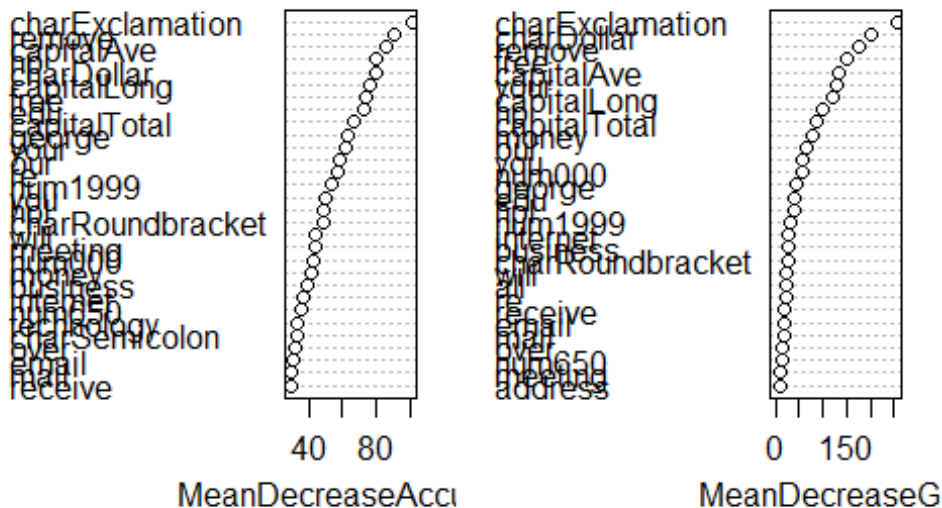


We see that random forests method out performs the other methods with an AUC of 98.7%

We train a random forest model with 2000 trees using the entire data set.

```
set.seed(5474)
spam.rf.full <- randomForest(type~., data = dat, importance = TRUE, proximity
                             = TRUE,
                             ntree = 2000)
varImpPlot(spam.rf.full, main = "Variable Importance")
```

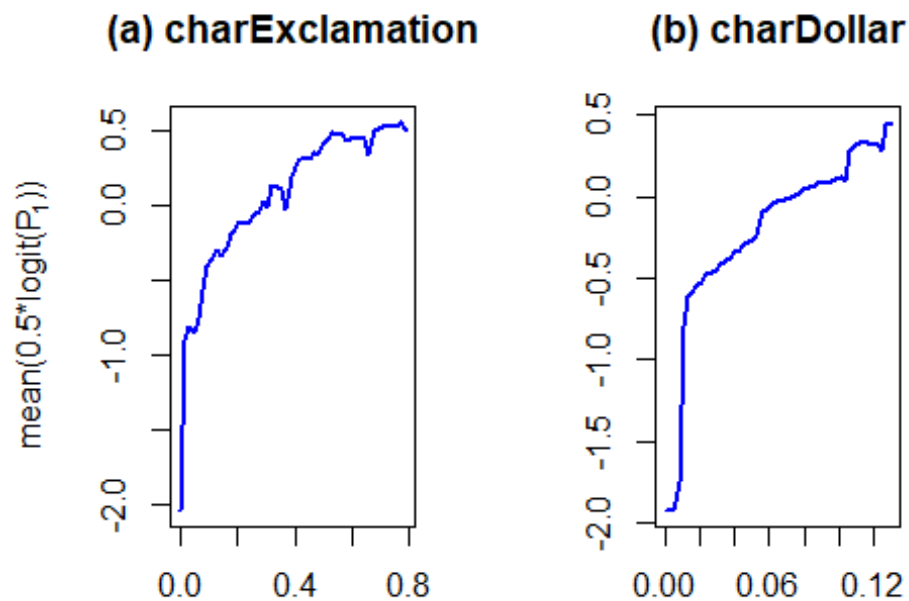
Variable Importance



Variable importance plot for the random forest fit with the full data set

From the variable importance plots we see that charExclamation and charDollar are the top two most important variables according to the mean decrease in Gini index. We next plot the partial dependence plots for these two variables.

```
library(interpretR, quietly = TRUE)
par(mfrow=c(1, 2))
parDepPlot(x.name = "charExclamation", spam.rf.full, data = dat,
  main = "(a) charExclamation", col = "blue", lwd = 2, xlab = "",
  yalb = "Partial Dependence")
parDepPlot(x.name = "charDollar", spam.rf.full, data = dat,
  main = "(b) charDollar", col = "blue", lwd = 2, xlab="", ylab="")
```

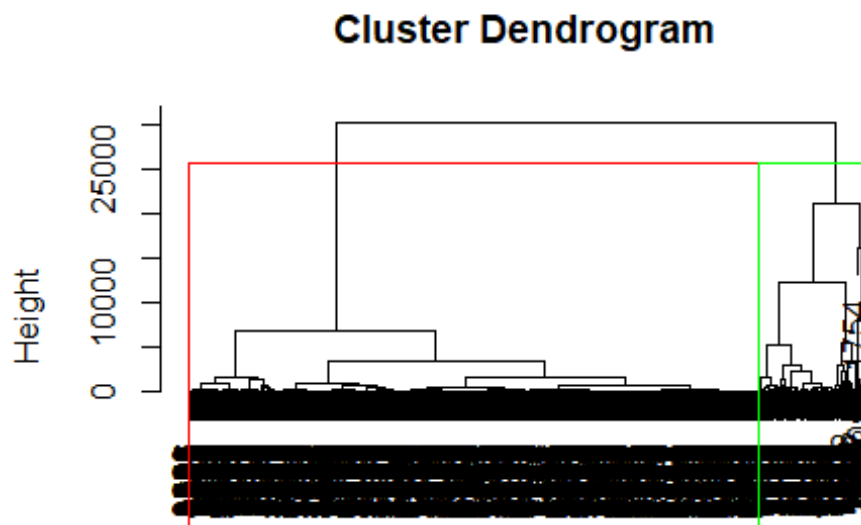


Partial dependence plots for top two important variables.

In the partial dependence plot, we see that as the number of exclamation marks (a) and dollar signs (b) increase, the more likely an email will be classified as a spam email.

We will be running hierarchical clustering technique, but this time with only the training data. We will also plot the tSNE denoting the new response variable and as well as the original.

```
distance.training <- get_dist(training[, -58], method = "euclidean")
hc.out.training <- hclust(distance.training, method = "ward.D2")
hc.clusters.training <- cutree(hc.out.training, k=2)
plot(hc.out.training)
rect.hclust(hc.out.training, k=2, border = c("red", "green"))
```

```
distance.training
hclust (*, "ward.D2")
```

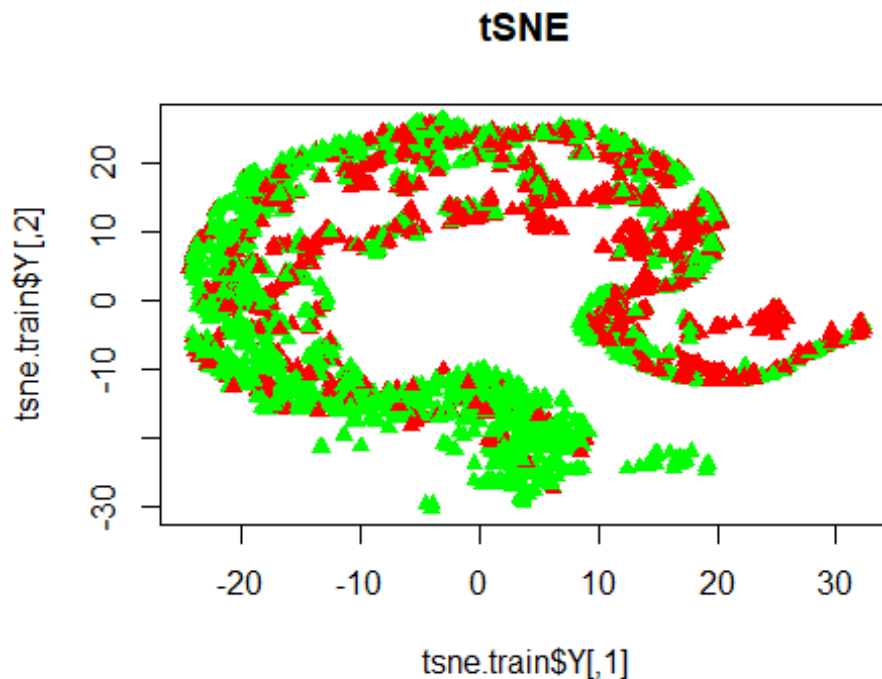
```
table(hc.clusters.training)/length(hc.clusters.training) * 100
```

```
## hc.clusters.training
##      1      2
## 15.41721 84.58279
```

Again, hierarchical clustering didn't perform well.

now for tSNE

```
new.train <- cbind(training, hc.clusters.training)
tsne.train <- Rtsne(new.train[, -c(58,59)], dims = 2,
                    perplexity = 43, verbose = FALSE,
                    check_duplicates = FALSE, max_iter = 500)
plot(tsne.train$Y, main = "tSNE", col = c("green", "red")[new.train[, 58]],
     pch = tsne.train$Y[new.train[, 59]])
```



The plot still hasn't

improved much.

```
n <- nrow(dat)
out <- NULL
for (k in 1:ncol(dat)){
  vname <- colnames(dat)[k]
  x <- as.vector(dat[,k])
  n1 <- sum(is.na(x), na.rm = TRUE)
  n2 <- sum(x=="NA", na.rm = TRUE)
  n3 <- sum(x=="'", na.rm = TRUE)
  n4 <- sum(x=="?", na.rm = TRUE)
  n5 <- sum(x=="*", na.rm = TRUE)
  n6 <- sum(x==".", na.rm = TRUE)
  nmiss <- n1 + n2 + n3 + n4 + n5 + n6
  ncomplete <- n - nmiss
  var.type <- typeof(x)
  if (var.type == "integer"){
    if (length(unique(x)) == 2){
      out <- rbind(out, c(col.number=k, vname=vname, mode="binary",
        n.levels=length(unique(x)), ncomplete = ncomplete,
        miss.prop=round(nmiss/n, digits = 4)))
    } else {
      out <- rbind(out, c(col.number=k, vname=vname, mode=typeof(x),
        n.levels=length(unique(x)), ncomplete=ncomplete,
        miss.prop=round(nmiss/n, digits = 4)))
    }
  } else {
```

```

    out <- rbind(out, c(col.number=k, vname=vname, mode=typeof(x),
                        n.levels=length(unique(x)), ncomplete=ncomplete,
                        miss.prop=round(nmiss/n, digits = 4)))
  }
}
out <- as.data.frame(out)
row.names(out) <- NULL
out

```

##	col.number	vname	mode	n.levels	ncomplete	miss.prop
## 1	1	make	double	142	4601	0
## 2	2	address	double	171	4601	0
## 3	3	all	double	214	4601	0
## 4	4	num3d	double	43	4601	0
## 5	5	our	double	255	4601	0
## 6	6	over	double	141	4601	0
## 7	7	remove	double	173	4601	0
## 8	8	internet	double	170	4601	0
## 9	9	order	double	144	4601	0
## 10	10	mail	double	245	4601	0
## 11	11	receive	double	113	4601	0
## 12	12	will	double	316	4601	0
## 13	13	people	double	158	4601	0
## 14	14	report	double	133	4601	0
## 15	15	addresses	double	118	4601	0
## 16	16	free	double	253	4601	0
## 17	17	business	double	197	4601	0
## 18	18	email	double	229	4601	0
## 19	19	you	double	575	4601	0
## 20	20	credit	double	148	4601	0
## 21	21	your	double	401	4601	0
## 22	22	font	double	99	4601	0
## 23	23	num000	double	164	4601	0
## 24	24	money	double	143	4601	0
## 25	25	hp	double	395	4601	0
## 26	26	hpl	double	281	4601	0
## 27	27	george	double	240	4601	0
## 28	28	num650	double	200	4601	0
## 29	29	lab	double	156	4601	0
## 30	30	labs	double	179	4601	0
## 31	31	telnet	double	128	4601	0
## 32	32	num857	double	106	4601	0
## 33	33	data	double	184	4601	0
## 34	34	num415	double	110	4601	0
## 35	35	num85	double	177	4601	0
## 36	36	technology	double	159	4601	0
## 37	37	num1999	double	188	4601	0
## 38	38	parts	double	53	4601	0
## 39	39	pm	double	163	4601	0

## 40	40	direct	double	125	4601	0
## 41	41	cs	double	108	4601	0
## 42	42	meeting	double	186	4601	0
## 43	43	original	double	136	4601	0
## 44	44	project	double	160	4601	0
## 45	45	re	double	230	4601	0
## 46	46	edu	double	227	4601	0
## 47	47	table	double	38	4601	0
## 48	48	conference	double	106	4601	0
## 49	49	charSemicolon	double	313	4601	0
## 50	50	charRoundbracket	double	641	4601	0
## 51	51	charSquarebracket	double	225	4601	0
## 52	52	charExclamation	double	964	4601	0
## 53	53	charDollar	double	504	4601	0
## 54	54	charHash	double	316	4601	0
## 55	55	capitalAve	double	2161	4601	0
## 56	56	capitalLong	double	271	4601	0
## 57	57	capitalTotal	double	919	4601	0
## 58	58	type character		2	4601	0