

Project-01: SEMMA with Regularized Logistic Regression

Ruben Solis
rsolis5@miners.utep.edu

Due 9/23/2019

Contents

1 Problem-1: Data Collection	2
2 Problem-2: Exploratory Data Analysis	2
2.1 Removing column ID	3
2.2 Inspecting the distinct values of each variable	3
2.3 Visualizing Missing values	4
2.4 Frequency Distribution of the Target variable 'Class'	8
2.5 Inspecting the outlying records	8
2.6 Association using χ^2 and Fisher test	9
2.6.1 Measure the association between Class and other variables	11
3 Problem-3: Data Partition	13
4 Problem 4:	13
4.1 Problem-4(a): Building a Logistic Rgression Model	13
4.2 Problem-4(b): Selecting the best tuning parameter	13
5 Problem-5	15
5.1 Final logistic model to the test data D3	15
5.2 ROC curve and AUC	16

1 Problem-1: Data Collection

We obtain the data directly in R using `mlbench` package and save the data in a CSV file. The code is as follows:

```
#install.packages("mlbench")
data(BreastCancer, package="mlbench")
data <- BreastCancer;
write.csv(data, file="BreastCancer.csv", row.names =FALSE)
```

2 Problem-2: Exploratory Data Analysis

To perform EDA, we first inspect the data types, quantity and percentage of zeros, infinite numbers, missing, and unique values of data using `df_status` function:

```
#install.packages("funModeling")
library("funModeling")
dim(data)
```

```
## [1] 699 11
```

```
head(data)
```

```
##      Id Cl.thickness Cell.size Cell.shape Marg.adhesion Epith.c.size
## 1 1000025          5         1         1           1           2
## 2 1002945          5         4         4           5           7
## 3 1015425          3         1         1           1           2
## 4 1016277          6         8         8           1           3
## 5 1017023          4         1         1           3           2
## 6 1017122          8        10        10           8           7
##  Bare.nuclei Bl.cromatin Normal.nucleoli Mitoses      Class
## 1           1           3           1         1    benign
## 2          10           3           2         1    benign
## 3           2           3           1         1    benign
## 4           4           3           7         1    benign
## 5           1           3           1         1    benign
## 6          10           9           7         1 malignant
```

```
#str(data)
```

```
df_status(data) # Getting the metrics about data types, zeros, infinite numbers, and missing values
```

```
##      variable q_zeros p_zeros q_na p_na q_inf p_inf      type
## 1      Id        0        0    0 0.00    0    0    character
## 2 Cl.thickness    0        0    0 0.00    0    0 ordered-factor
## 3   Cell.size     0        0    0 0.00    0    0 ordered-factor
## 4   Cell.shape    0        0    0 0.00    0    0 ordered-factor
## 5 Marg.adhesion    0        0    0 0.00    0    0 ordered-factor
## 6 Epith.c.size     0        0    0 0.00    0    0 ordered-factor
## 7  Bare.nuclei     0        0   16 2.29    0    0      factor
```

```
## 8      Bl.cromatin      0      0      0 0.00      0      0      factor
## 9 Normal.nucleoli      0      0      0 0.00      0      0      factor
## 10      Mitoses        0      0      0 0.00      0      0      factor
## 11      Class          0      0      0 0.00      0      0      factor
##      unique
## 1      645
## 2      10
## 3      10
## 4      10
## 5      10
## 6      10
## 7      10
## 8      10
## 9      10
## 10     9
## 11     2
```

2.1 Removing column ID

We prepare the data by removing column ID, Since it will not affect in our analysis.

```
dat <- data[, -1]
dim(dat)
```

```
## [1] 699 10
```

2.2 Inspecting the distinct values of each variable

We inspect the distinct values of each variable of data. The following code shows that the target Class has two categorical levels: 458 benign and 241 malignant. The other variable is also shown 10 levels with number of values.

```
for (j in 1:NCOL(dat)){
  print(colnames(dat)[j])
  print(table(dat[,j], useNA="ifany"))
}
```


```
## [1] "Cl.thickness"
##
##  1  2  3  4  5  6  7  8  9 10
## 145 50 108 80 130 34 23 46 14 69
## [1] "Cell.size"
##
##  1  2  3  4  5  6  7  8  9 10
## 384 45 52 40 30 27 19 29 6 67
## [1] "Cell.shape"
##
##  1  2  3  4  5  6  7  8  9 10
```

```
## 353  59  56  44  34  30  30  28   7  58
## [1] "Marg.adhesion"
##
##   1   2   3   4   5   6   7   8   9  10
## 407  58  58  33  23  22  13  25   5  55
## [1] "Epith.c.size"
##
##   1   2   3   4   5   6   7   8   9  10
##  47 386  72  48  39  41  12  21   2  31
## [1] "Bare.nuclei"
##
##   1   2   3   4   5   6   7   8   9  10 <NA>
## 402  30  28  19  30   4   8  21   9 132  16
## [1] "Bl.cromatin"
##
##   1   2   3   4   5   6   7   8   9  10
## 152 166 165  40  34  10  73  28  11  20
## [1] "Normal.nucleoli"
##
##   1   2   3   4   5   6   7   8   9  10
## 443  36  44  18  19  22  16  24  16  61
## [1] "Mitoses"
##
##   1   2   3   4   5   6   7   8  10
## 579  35  33  12   6   3   9   8  14
## [1] "Class"
##
##      benign malignant
##      458         241
```


2.3 Visualizing Missing values

We now visualize the quantity, percentage of missing values for each variables. We see that only the Bare.nuclei variable has 16 NA values, which is 2.29 percent of the data. We also visualize the the missing values for categorical variables and then we omit the NA data.

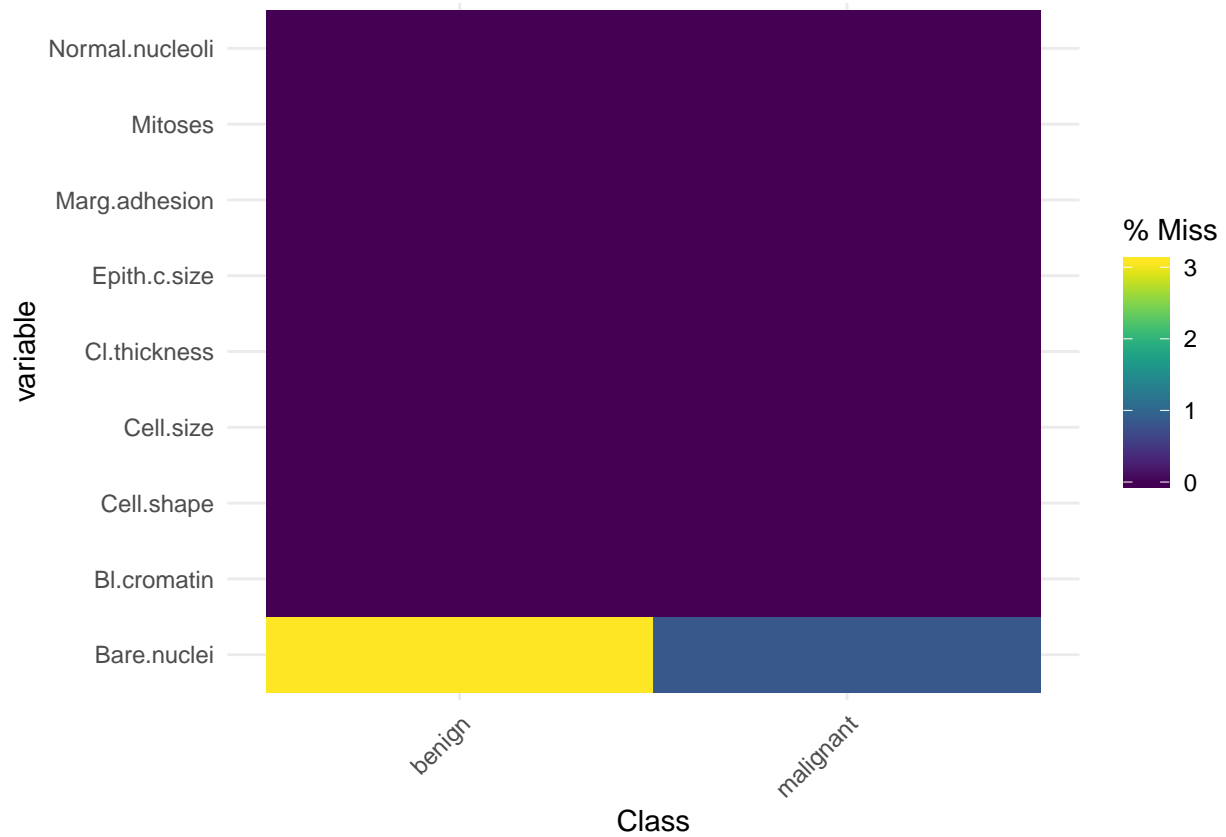
```
#install.packages("naniar")
library(naniar)
vis_miss(dat)
```



```
gg_miss_var(dat)
```



```
##Visulaizing the missing values for target Class of each variables  
gg_miss_fct(x = dat, fct = Class)
```




```
# remove rows containing missing values
dat <- na.omit(dat)
dim(dat)
```

```
## [1] 683 10
```

2.4 Frequency Distribution of the Target variable 'Class'

The following code shows the frequency, percentage, cumulative percentage of the target variable 'Class'.

```
freq(dat$Class)
```



```
##          var frequency percentage cumulative_perc
## 1    benign      444      65.01           65.01
## 2 malignant      239      34.99          100.00
```

To analyze the data, we make numeric values 1 and 0 for categorical variable Class.

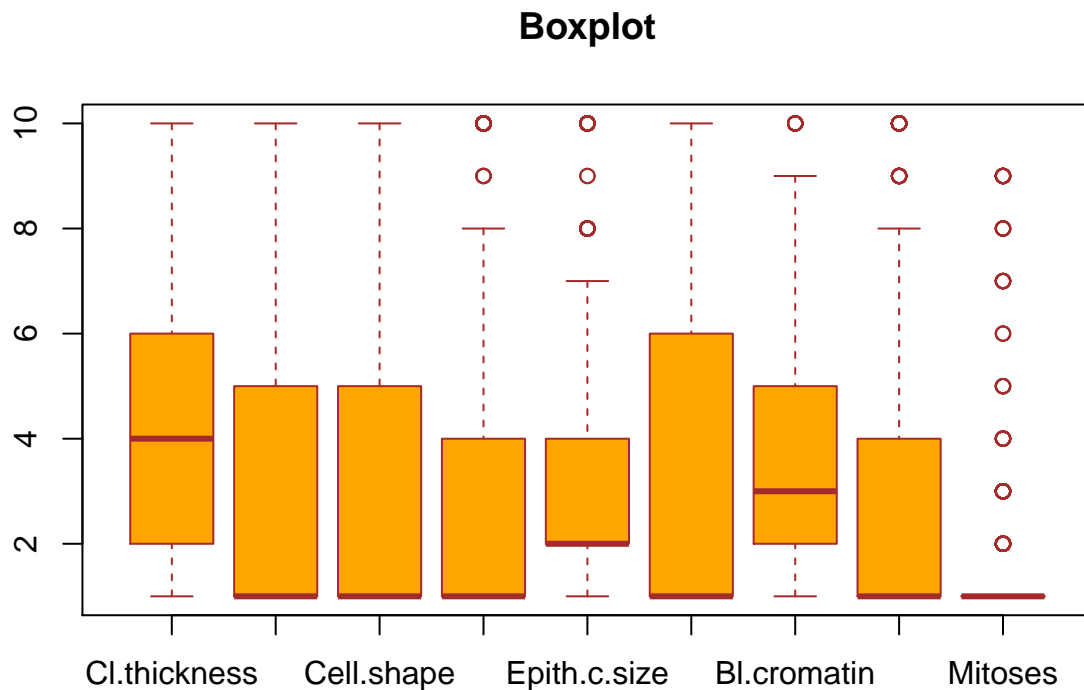
```
dat$y <- ifelse(dat$Class=="benign", 1, 0)
dim(dat)
```

```
## [1] 683  11
```

2.5 Inspecting the outlying records

We now plot a Box plot to inspect the outlying records for each variable. As we see in the graph, Marg adhesion, Epith.c.size, B1 cromatin, Normal Nucleoli has few outlying value. Mitoses variable has most of the outlying values.


```
data_box <- dat[, -11:-10] # Removing Class variable (Bcz of many variables)
boxplot(data_box, main = "Boxplot", horizontal = F, col="orange",
        border="brown")
```



To analyze the data, we change the character variable of into numeric variable.

```
dat <- dat[, c(1:9, 11)]
dat <- apply(dat, 2, FUN=function(x) {as.numeric(as.character(x))})
dat <- na.omit(dat)
dat <- as.data.frame(dat)
```

2.6 Association using χ^2 and Fisher test

To check the association between Class and other attributes, we use χ^2 and fisher test. Our hypotheses are as follows:

H_0 : The two variables are independent,

H_1 : The two variables are dependent.

since the p-value is less than the significance level (0.05) for all cases (between class and other attributes), we reject the null hypothesis and conclude that the Class and other attribute are dependent to each other. The code and result are as follows:

```

#####
# Chi-sq and Fisfter test for Class and Other varibales
#####
dat1 <- dat[, -11]
chitest <- matrix (0, 9, 4)
ftest <- matrix (0, 9, 2)

for (j in 1: (ncol(dat1)-1)){
  testor <- table(as.vector(dat1 [, ncol(dat1)]), as.numeric(dat1[, j]))
  chi2 <- chisq.test(testor, correct=FALSE)
  chitest[j, ] <- c(colnames(dat1)[j], round(chi2$statistic, digits = 2), chi2$p.value, chi2$par)
  s = fisher.test(testor, simulate.p.value = TRUE, B=1e5)
  ftest[j, ] <- c(colnames(dat1)[j], s$p.value)
}

colnames(chitest) <- c("Names", "Statistics", "p-values", "D.Freedom")
colnames(ftest) <- c("Names", "p-values")
names(dimnames(chitest)) <- list("", "Association among Class and other predictors Using Chi test")
names(dimnames(ftest)) <- list("", "Association among Class and other predictors Using Fisher test")
chitest

```

```

##          Association among Class and other predictors Using Chi test
##
##      Names      Statistics p-values      D.Freedom
## [1,] "Cl.thickness" "378.08"  "6.47143951518931e-76" "9"
## [2,] "Cell.size"    "539.79"  "1.71638971098766e-110" "9"
## [3,] "Cell.shape"   "523.07"  "6.57844762931427e-107" "9"
## [4,] "Marg.adhesion" "390.06"  "1.80795747026517e-78"  "9"
## [5,] "Epith.c.size" "447.86"  "8.21759531792845e-91"  "9"
## [6,] "Bare.nuclei"  "489.01"  "1.2957665166586e-99"   "9"
## [7,] "Bl.cromatin"  "453.21"  "5.90593696241214e-92"  "9"
## [8,] "Normal.nucleoli" "416.63" "3.86380729078817e-84"  "9"
## [9,] "Mitoses"      "191.97"  "3.13852341562463e-37"  "8"

```

```
ftest
```

```

##          Association among Class and other predictors Using Fisher test
##
##      Names      p-values
## [1,] "Cl.thickness" "9.99990000099999e-06"
## [2,] "Cell.size"    "9.99990000099999e-06"
## [3,] "Cell.shape"   "9.99990000099999e-06"
## [4,] "Marg.adhesion" "9.99990000099999e-06"
## [5,] "Epith.c.size" "9.99990000099999e-06"
## [6,] "Bare.nuclei"  "9.99990000099999e-06"
## [7,] "Bl.cromatin"  "9.99990000099999e-06"
## [8,] "Normal.nucleoli" "9.99990000099999e-06"
## [9,] "Mitoses"      "9.99990000099999e-06"

```

2.6.1 Measure the association between Class and other variables


In this subsection, we study the association plot (given below) where the diagonal element K refers to number of unique levels for each variable. This measure of association indicates the strength of the relationship, whether, weak or strong. The off-diagonal elements contain the forward and backward τ measures for each variable pair. Specifically, the numerical values appearing in each row represent the association measure $\tau(x, y)$ from the variable xx indicated in the row name to the variable yy indicated in the column name.

For example, the variable `Cell.size` is almost perfectly predictable (i.e. $\tau(x, y) = 0.78$) from `Class` and this forward association is quite strong. The forward association suggest that $x=\text{Cell.size}$ is highly predictive of $y=\text{Class}$. It indicates that if we know a `Cell.size`, then we can easily predict its `Class`.

On the contrary, the reverse association $y=\text{class}$ and $x=\text{Cell.size}$ (i.e. $\tau(y, x) = 0.27$); is a strong association and indicates that if we know the `Class` then its easy to predict its `Cell.size`.

From chi-squared and Fisher significance test, we have found `Normal.Neocleoli` and `Cell.thickness` are dependent to each other. But forward and reverse association plot suggest that $x=\text{Normal.neocleuli shape}$ is weakly associated to $y=\text{Cell.thickness}$ (i.e. $\tau(x, y) = 0.17$) and (i.e. $\tau(y, x) = 0.060$). So we conclude that although these two variables are significant but their association is weak; i.e. it will be difficult to predict one from another.

```
#install.packages('GoodmanKruskal')
library(GoodmanKruskal)
varset1<- c("Cl.thickness", "Cell.size", "Cell.shape", "Marg.adhesion", "Epith.c.size", "Bar")
associate1<- subset(data, select = varset1)
GKmatrix1<- GKtauDataframe(associate1)
plot(GKmatrix1, corrColors = "blue")
```



3 Problem-3: Data Partition

In this section, we partition the data into three parts, the training data D1, the validation data D2, and the test data D3, with a ratio of 2 : 1 : 1.

```
set.seed(123)
n <- nrow(dat)
id.split <- sample(x=1:3, size = n, replace =TRUE, prob=c(0.5, 0.25, 0.25))
dat.train <- dat[id.split ==1, ]
dat.valid <- dat[id.split ==2, ]
dat.test <- dat[id.split == 3, ]
```

4 Problem 4:

4.1 Problem-4(a): Building a Logistic Rgression Model

Here, we fit the regularized logistic regression using the training data \mathcal{D}_∞ with the Lasso model. In the glmnet function, the family argument, specify that we want a “binomial” model which tells glmnet() to fit a logistic function to the data.

```
#install.packages("glmnet")
library(glmnet)
formula0 <- y~Cl.thickness + Cell.size + Cell.shape + Marg.adhesion + Epith.c.size + Bare.nucl
X <- model.matrix (as.formula(formula0), data = dat.train)
y <- dat.train$y
```

4.2 Problem-4(b): Selecting the best tuning parameter

Next, we would like to see how the model is doing when predicting Class(y) on data. we use `pred` function in the form of $P(y = 1|X)$ using parameter `type='response'` which tells predict to return probabilities. The decision boundary will be 0.5. If $P(y = 1|X) > 0.5$ then $y = 1$ (benign) otherwise $y = 0$ (malignant). Therefore, we used misclassification rate and the mean square error (MSE) for the predicted probabilities.

Now, we select the best tuning parameter using the validation data \mathcal{D}_ϵ and choosing the minimum mean squared error (MSE).

```
X.valid <- model.matrix (as.formula(formula0), data = dat.valid)
y.valid <- dat.valid$y

Lambda <- seq(0.0001, 0.5, length.out = 200)
L <- length(Lambda)
OUT <- matrix (0, L, 3)
for (i in 1:L){
  fit <- glmnet(x=X, y=y, family ="binomial", alpha =1, #lasso
               lambda=Lambda[i], standardize=T, thresh = 1e-07, maxit=1000)
  pred <- predict(fit, newx=X.valid, s=Lambda[i], type="response")
```


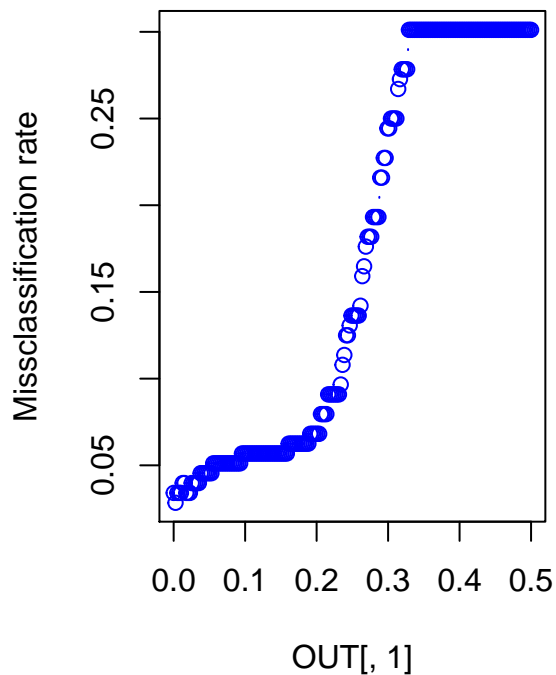
```

miss.rate <- mean(y.valid != (pred > 0.5))
mse <- mean((y.valid - pred)^2)
OUT[i, ] <- c(Lambda[i], miss.rate, mse)
}
head(OUT)

##           [,1]      [,2]      [,3]
## [1,] 0.000100000 0.03409091 0.02181988
## [2,] 0.002612060 0.02840909 0.02123402
## [3,] 0.005124121 0.03409091 0.02156797
## [4,] 0.007636181 0.03409091 0.02204101
## [5,] 0.010148241 0.03409091 0.02254983
## [6,] 0.012660302 0.03977273 0.02308470

par(mfrow = c(1,2))
plot(OUT[, 1], OUT[,2], type = "b", col = "blue", ylab = "Missclassification rate")
plot(OUT[, 1], OUT[,3], type = "b", col = "red", ylab = "MSE")

```



```

(lambda.best <- OUT[which.min(OUT[, 3]), 1])

```

```
## [1] 0.00261206
```

```
(miss.rate_Lambda <- OUT[which.min(OUT[, 3]), 2])
```

```
## [1] 0.02840909
```

The best fitted lambda is 0.00261206, where the corresponding misclassification rate is 0.02840909. So we can conclude that the accuracy on this model is good.

\subsection{Problem-4(c): Final ‘best’ model by pooling D1 and D2 We then present the final “best” model fit by pooling \mathcal{D}_∞ and D2 together. Using the beta from `fit.best` model, I see the coefficients of the model and select the important predictors.

```
X.12 = rbind(X, X.valid)
y.12 = c(y, y.valid)
fit.best <- glmnet (x=X.12, y=y.12, family ="binomial", alpha=1, #LASSO
                    lambda = lambda.best, standardize = T, thresh = 1e-07, maxit=1000)
names(fit.best)
```

```
## [1] "a0"      "beta"    "df"      "dim"     "lambda"
## [6] "dev.ratio" "nulldev" "npasses" "jerr"    "offset"
## [11] "classnames" "call"    "nobs"
```

```
fit.best$beta # Finding important variables.
```

```
## 10 x 1 sparse Matrix of class "dgCMatrix"
##                               s0
## (Intercept)                   .
## Cl.thickness    -0.4891238
## Cell.size       .
## Cell.shape      -0.2656879
## Marg.adhesion   -0.3596817
## Epith.c.size    -0.2128013
## Bare.nuclei     -0.2988293
## Bl.cromatin     -0.3582619
## Normal.nucleoli -0.1435005
## Mitoses         -0.3063777
```

Since we do not get any coefficients for `Cell.size`, it is not an important predictor for the predictive model. The important predictors are `Cl.thickness`, `Cell.shape`, `Marg.adhesion`, `Epith.c.size`, `Bare.nuclei`, `Bl.cromatin`, `Normal.nucleoli`, and `Mitoses`.

5 Problem-5

5.1 Final logistic model to the test data D3

Using test data D3, we apply the final logistic model.

```
X.test <- model.matrix(object=~Cl.thickness + Cell.size +Cell.shape + Marg.adhesion + Epith.c.
pred <- predict(fit.best, newx = X.test, s =lambda.best, type="response")
dim(pred)
```

```
## [1] 158    1
```

5.2 ROC curve and AUC

ROC suggests the accuracy of a classification model at a threshold value. It determines the model's accuracy using Area Under Curve (AUC). The AUC also referred to as index of accuracy (A) or concordant index (ci), represents the performance of the ROC curve. The idea is that higher the area, better the model.


```
library(cvAUC)
yobs <- dat.test$y
AUC <- ci.cvAUC(predictions = pred, labels = yobs, folds=1:NROW(dat.test), confidence = 0.95)
AUC

## $cvAUC
## [1] 0.9975066
##
## $se
## [1] 0.001775487
##
## $ci
## [1] 0.9940268 1.0000000
##
## $confidence
## [1] 0.95

auc.ci <- round(AUC$ci, digits = 3)

library(verification)
mod.glm <- verify(obs = yobs, pred = pred)

## If baseline is not included, baseline values will be calculated from the sample obs.
roc.plot(mod.glm, plot.thres=NULL)
text(x=0.7, y=0.2, paste("Area under ROC = ", round(AUC$cvAUC, digits = 3), "with 95% CI (",
                        auc.ci[1], ",", auc.ci[2], ").", sep = " "), col="blue", cex=1.2)
```

ROC is plotted between True Positive Rate (Y axis) and False Positive Rate (X axis). In this plot, our aim is to push the curve (shown below) toward 1 (left corner) and maximize the area under curve. The diagonal line represents the ROC curve at 0.5 threshold. Since AUC is 0.99 (closer to 1) and the curve almost approaches to 1, we can say that the model have good predictive ability.